

# Neural Networks

**Neural networks** are a class of **machine learning** models inspired by the **brain**. They are non-linear models that can be used for **regression** and **classification**.

## Definition

- Inspired by biological **neurons**;
- An **artificial neuron** is a **function**  $f : \mathbb{R}^d \rightarrow \mathbb{R}$ ;
  - It receives a **vector** of **inputs**  $x \in \mathbb{R}^d$ , **weights**  $w \in \mathbb{R}^d$  and a **bias**  $b \in \mathbb{R}$ :

$$z(x) = w^T x + b = \sum_{i=1}^d w_i x_i + b$$

The **activation function**  $g$  is a function that **transforms** the **output** of the **neuron**:

$$h(x) = g(z(x)) = g(w^T x + b), \text{ where } g : \mathbb{R} \rightarrow \mathbb{R}$$

---

## Activation Functions

The typical **activation functions** are:

- **Linear**:  $g(z) = z$ ;
  - No **squashing**;
- **Sigmoid**:  $g(z) = \sigma(z) = \frac{1}{1+e^{-z}}$ ;
  - Squashes  $z$  to  $[0, 1]$ ;
  - Output can be interpreted as a **probability**;
  - Positive, bounded and strictly increasing;

- **Hyperbolic tangent:**  $g(z) = \tanh(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$ ;
  - Squashes  $z$  to  $[-1, 1]$ ;
  - Bounded and strictly increasing;
- **Rectified Linear Unit (ReLU):**  $g(z) = \text{relu}(z) = \max(0, z)$ ;
  - Non-negative, increasing but **not upper bounded**;
  - Not differentiable at  $z = 0$ ;
  - Leads to neurons with **sparse activities**-

Later, we will see:

- **Softmax:**  $g(z) = \frac{e^{z_i}}{\sum_{j=1}^d e^{z_j}}$ ;
- **Sparsemax:**  $g(z) = \operatorname{argmin}_{p \in \Delta_d} \|p - z\|_2^2$ ;
- **Max-pooling:**  $g(z) = \max(z)$ .

## Feedforward Neural Networks

- To solve non-linear problems, we can **stack** several **neurons**;
- **Multi-layer neural networks:** use intermediate layers between the input and the output layers;
- Each hidden layer computes a representation of the input and propagates it to the next layer - **feedforward neural networks**.

### Single Hidden Layer

- Consider a task that involves several **inputs**  $x \in \mathbb{R}^d$  and a **single output**  $y \in [0, 1]$ ;
- Include an **intermediate layer** of  $K$  **hidden units** ( $h \in \mathbb{R}^K$ ) between the input and the output layers;
- **Hidden layer pre-activation:**  $z(x) = W^{(1)}x + b^{(1)}$ , with  $W^{(1)} \in \mathbb{R}^{K \times d}$  and  $b^{(1)} \in \mathbb{R}^K$ ;
- **Hidden layer activation:**  $h(x) = g(z(x))$ , with  $g$  being the **activation function**  $g: \mathbb{R}^K \rightarrow \mathbb{R}^K$ ;
- **Output layer activation:**  $y(x) = \sigma(w^{(2)}h(x) + b^{(2)})$ , with  $w^{(2)} \in \mathbb{R}^K$  and  $b^{(2)} \in \mathbb{R}$ .

## Multiple Classes

- **Multiple output units**, one for each class;
- Each output estimates the conditional probability of the input belonging to that class - **softmax** activation function:

$$o(z) = \text{softmax}(z) = \begin{bmatrix} \frac{e^{z_1}}{\sum_{j=1}^d e^{z_j}} \\ \vdots \\ \frac{e^{z_d}}{\sum_{j=1}^d e^{z_j}} \end{bmatrix}$$

- **Hidden layer pre-activation:**  $z^{(l)}(x) = W^{(l)}h^{(l-1)}(x) + b^{(l)}$ , with  $W^{(l)} \in \mathbb{R}^{K^{(l)} \times K^{(l-1)}}$  and  $b^{(l)} \in \mathbb{R}^{K^{(l)}}$ ;
  - **Hidden layer activation:**  $h^{(l)}(x) = g(z^{(l)}(x))$ , with  $g$  being the **activation function**  $g: \mathbb{R}^{K^{(l)}} \rightarrow \mathbb{R}^{K^{(l)}}$ ;
  - **Output layer activation:**  $y(x) = \text{softmax}(w^{(L)}h^{(L-1)}(x) + b^{(L)})$ , with  $w^{(L)} \in \mathbb{R}^{K^{(L)}} \rightarrow \mathbb{R}$ .
- 

## Universal Approximation Theorem

*A Neural Network with a **single hidden layer** and a **linear output layer** can **approximate any continuous function to arbitrary accuracy** if the **hidden layer has enough units**. - Cybenko, 1989*

- **Deeper networks** (with more hidden layers) can **approximate functions more efficiently**;
  - The number of linear regions carved out by a deep neural network with  $D$  inputs, depth  $L$  and  $K$  hidden units per layer is  $O((\frac{K}{D})^{D(L-1)} K^D)$ ;
  - So neural networks can **approximate any function** with a **single hidden layer**: we only need to find the **right parameters** - **training**.
- 
- 

## Training Neural Networks

Training a neural network means **finding the right parameters** for the **weights** and **biases** of the **neurons**, learning them from data - samples of **inputs** and **outputs**.

## Empirical Risk Minimization

- Goal: choose parameters  $\theta := (W^{(l)}, b^{(l)})_{l=1}^{L+1}$  that **minimize the empirical risk**:

$$\mathcal{L}(\theta) := \lambda \Omega(\theta) + \frac{1}{N} \sum_{n=1}^N L(f(x_n; \theta), y_n)$$

- $x_i, y_{i=1}^N$  is the **training set**;
- $L(f(x_i; \theta), y_i)$  is the **loss function**;
- $\omega(\theta)$  is the **regularization term**;
- $\lambda$  is the **regularization constant**.
- **Gradient descent** is **too slow**, because it requires a full pass over the data to update the weights;
- **Stochastic gradient descent** is **faster**, because it updates the weights after each sample;
- **Mini-batch stochastic gradient descent** is **even faster**, because it updates the weights after each mini-batch of samples;
  - A **mini-batch** is a **subset** of the **training set**,  $j_1, \dots, j_B$ , with  $B$  samples,  $B \ll N$ :

$$\nabla_{\theta} \mathcal{L}(\theta) = \lambda \nabla_{\theta} \omega(\theta) + \frac{1}{B} \sum_{i=1}^B \nabla_{\theta} L(f(x_{j_i}; \theta), y_{j_i})$$

The weights are updated as follows:

$$\theta^{(t+1)} = \theta^{(t)} - \eta \frac{1}{B} \sum_{i=1}^B \nabla_{\theta} L_{j_i}(\theta^{(t)})$$

- **Loss function**  $L$  should match as possible what we want to **optimize**;
  - Should be well-behaved - continuous and smooth;
  - **Squared loss** is a good choice for **regression**;
  - **Cross-entropy** loss is a good choice for **multi-class classification**:  
 $L(f(x; \theta), y) = -\log(\text{softmax}(f(x; \theta))_y)$ ;
  - **Sparsemax** loss is a good choice for **multi-class and multi-label classification**.

## Backpropagation

We need to find a **procedure** to **compute the gradients** of the **loss function** with respect to the **weights** and **biases** of the **neurons**:  $\nabla_{\theta} L(f(x; \theta), y)$ .

- The **gradient backpropagation algorithm** is a **recursive algorithm** that computes the **gradients** of the **loss function** with respect to the **weights** and **biases** of the **neurons**;
- It is based on the **chain rule** of **calculus**:

$$h(x) = f(g(x)) \Rightarrow \frac{dh}{dx} = \frac{df}{dg} \frac{dg}{dx}$$

---

## Automatic Differentiation

- **Automatic differentiation** is a **technique** for **computing derivatives** of **functions**;
- Forward propagation can be represented as a **computation graph** - a **directed acyclic graph** (DAG) that represents the **computation** of the **function**;
  - Each box can be an object with a **fprop** method that computes the **forward pass**;
  - Calling the **fprop** method of each box in the **topological order** of the graph computes the **forward pass**;
- **Backpropagation** is also implemented as a **computation graph** - a **directed acyclic graph** (DAG) that represents the **computation** of the **gradients**;
  - Each box can be an object with a **bprop** method that computes the **loss gradient** w.r.t. its parents, given the **loss gradient** w.r.t. to the output of the box;
  - Calling the **bprop** method of each box in the **reverse topological order** of the graph computes the **backward pass**.

There are several **Autodiff** strategies:

- **Symbol-to-Number Differentiation**
  - Take a computational graph and numerical inputs;
  - Returns a set of numerical outputs describing the gradient at those inputs;
  - **Advantage**: simpler to implement and debug;

- **Disadvantage:** only works for first-order derivatives;
- **Example:** Caffe, Torch, PyTorch, ...

- **Symbol-to-Symbol Differentiation**

- Take a computational graph and add additional nodes to the graph that provide a symbolic description of the gradient;
- **Advantage:** works for higher-order derivatives;
- **Disadvantage:** more complex to implement and debug;
- **Example:** Theano, TensorFlow, ...

## Regularization

Recall the **empirical risk minimization** problem:

$$\mathcal{L}(\theta) := \lambda \Omega(\theta) + \frac{1}{N} \sum_{n=1}^N L(f(x_i; \theta), y_i)$$

- It remains to define the **regularizer**  $\Omega(\theta)$  and its **gradient**  $\nabla_{\theta} \Omega(\theta)$ ;

### $\Downarrow_2$ Regularization

- Only the **weights** are regularized:  $\Omega(\theta) = \sum_{l=1}^{L+1} \|W^{(l)}\|^2$ ;
- Equivalent to **Gaussian prior** on the weights;
- **Gradient of the regularizer** w.r.t. the **weights**:  $\nabla_{W^{(l)}} \Omega(\theta) = W^{(l)}$ ;
- **Weight decay effect:** the **weights** are **shrunk** towards **zero**:  $W^{(l)} \leftarrow (1 - \eta\lambda)W^{(l)}$ ;

### $\Downarrow_1$ Regularization

- Only the **weights** are regularized:  $\Omega(\theta) = \sum_l \|W^{(l)}\|_1 = \sum_l \sum_{i,j} |W_{ij}^{(l)}|$ ;
- Equivalent to **Laplace prior** on the weights;
- **Gradient of the regularizer** w.r.t. the **weights**:  $\nabla_{W^{(l)}} \Omega(\theta) = \text{sign}(W^{(l)})$ ;
- Promotes **sparsity** in the weights: **many weights** are **zeroed out**.

## Dropout Regularization

- During training, **randomly drop** some of the **neurons** in the **hidden layers**;
  - Each hidden unit output is set to zero with probability  $p$  - this prevents hidden units to **co-adapt** to each other, forcing them to be more generally useful;
  - At test time, keep all units with the **outputs multiplied by  $1 - p$** ;
  - Usually implemented using **random binary masks**;
  - The hidden layer activations becomes:  $h^{(\dagger)}(x) = g(z^{(\dagger)}(x)) \odot m^{(\dagger)}$ , where  $m^{(\dagger)} \in \{0, 1\}^{K^{(\dagger)}}$  is a **random binary mask** with  $p$  probability of being 1;
- 
- 

## Tricks of the Trade

### Initialization

- **Biases**: set to zero;
- **Weights**:
  - Cannot be zero with **tanh** activation function;
  - Cannot be all the same value - use **random initialization - Gaussian** or **uniform**;
  - For **ReLU**, the mean should be a small positive number;
  - Variance cannot be too high;

### Training, Validation and Test Sets

- **Training set**: used to **train** the model;
- **Validation set**: used to **tune** the **hyperparameters** (e.g. learning rate, regularization constant, etc.);
  - **Grid search** specify a set of values to test for each hyperparameter, and try all combinations;
  - **Random search** specify a distribution for each hyperparameter, and sample from it;
  - **Bayesian optimization** specify a prior distribution for each hyperparameter, and update it after each experiment;

- **Test set:** used to **evaluate** the **final model**.

**Early stopping** is a **regularization technique** that stops training when the **validation error** starts to **increase**, in order to prevent **overfitting**.

### Input Normalization

- Subtract the mean and divide by the standard deviation;
- It makes each input dimension have **zero mean** and **unit variance**;
- It **speeds up** the **training**;
- Does not work for **sparse data**;

### Decaying the Learning Rate

- **Learning rate**  $\eta$  is a **hyperparameter** that controls the **step size** in the **gradient descent** algorithm;
- In SGD, as we get closer to the minimum, we want to **reduce the step size**:
  - Start with a large learning rate (e.g. 0.1);
  - Keep it fixed while the **validation error** is **decreasing**;
  - Divide by 2 and repeat.

### Mini-Batches

- Instead of updating after a single sample, update after a **mini-batch** of samples (e.g. 50 – 200 samples), and compute the average gradient for the entire mini-batch;
- Less noisy than SGD;
- Can leverage matrix-matrix computations.

### Gradient Checking

- If the training loss is **not decreasing**, there might be a **bug** in the **gradient computation**;
- To debug, we can compute the **numerical gradient** and compare it with the **analytical gradient**:

$$\frac{\partial f}{\partial x} \approx \frac{f(x + \epsilon) - f(x - \epsilon)}{2\epsilon}$$


---



---



## Better Optimization

There are several improvements to basic **gradient descent** and **stochastic gradient descent**:

- **Momentum**;
- **Adaptive gradient** (AdaGrad);
- **Root mean square propagation** (RMSProp);
- **Adaptive moment estimation** (Adam).

### Momentum

- **Momentum** is a **technique** that **accelerates gradient descent** in the **relevant direction** and **dampens** oscillations;
- It means: remember the **previous gradients** and use them to **update** the **current gradient**:  $\theta_t = \theta_{t-1} - \alpha_t g(\theta_{t-1}) + \gamma_t(\theta_{t-1} - \theta_{t-2})$ ;  
–  $g(\theta_t)$  is the **gradient estimate** at time  $t$ ;
- **Advantages**: reduces the update in irrelevant directions and accelerates the update in relevant directions.

### Adaptive Gradient (AdaGrad)

- **AdaGrad** is a **technique** that **adapts the learning rate** to the **parameters**, performing **smaller updates** for **frequent** parameters and **larger updates** for **infrequent** parameters;
- Scale the update of each component ( $\epsilon$  for numerical stability):  $\theta_{j,t} = \theta_{j,t-1} - \frac{\alpha}{\sqrt{G_{j,t-1} + \epsilon}} g_j(\theta_{t-1})$ ;  
–  $G_{j,t}$  is the **sum of squares of the gradients** w.r.t.  $\theta_j$  up to time  $t$ :  $G_{j,t} = \sum_{i=1}^t g_{j,i}^2$ ;
- **Advantages**: robust to choice of  $\alpha$  and **learning rate decay**;
- **Disadvantages**: step size vanishes, because  $G_{j,t}$  is monotonically increasing.

### Root Mean Square Propagation (RMSProp)

- **RMSProp** addresses the **vanishing learning rate** issue;
- Same scaled update for each component:  $\theta_{j,t} = \theta_{j,t-1} - \frac{\alpha}{\sqrt{G_{j,t-1} + \epsilon}} g_j(\theta_{t-1})$ ;
- Use a **forgetting factor**  $\gamma$  to **decay the sum of squares of the gradients** (typically  $\gamma = 0.9$ ):  $G_{j,t} = \gamma G_{j,t-1} + (1 - \gamma) g_{j,t}^2$  - now the **sum of squares of the gradients** is **decaying**.

## Adaptive Moment Estimation (Adam)

- **Adam** is a **combination** of **momentum** and **RMSProp**;
- Separate moving averages of gradient and squared gradient;
- **Bias correction** is used to initialize the moving averages to zero;
- **Hyperparameters**:  $\alpha$ ,  $\beta_1$ ,  $\beta_2$ ,  $\epsilon$ ;
- **Advantages**: computationally efficient, low memory requirements, well suited for problems with large datasets and/or parameters;
- **Disadvantages**: possible convergence issues and noisy gradient estimates.