

Representation Learning

Representation learning is a set of techniques that **allow a system to automatically discover the representations needed for feature detection or classification** from raw data.

- A key feature of NNs is their ability to **learn representations** of the **data** $\phi(x)$;
- Standard linear models require **hand-crafted features**;
- **Representations** are useful for several reasons:
 - Can make models more **expressive and accurate**;
 - They may allow transferring representations from one task to another.

Hierarchical Compositionality

- Deep NNs learn **coarse-to-fine** representation layers;
- **Hierarchical compositionality** is the idea that **complex concepts** are composed of **simpler ones**;
- Layer closer to inputs learn **simple concepts** - edges, corners, etc.;
- Layer closer to outputs learn **more abstract representations** - shapes, forms, objects, etc.

But now some questions arise:

- How can a NN so effectively **represent** and **manipulate** knowledge, if it has only a few hidden units?
 - What is each hidden unit actually **representing**?
 - How can a NN **generalize** to objects that it has not seen before?
-

Distributed Representations

- **Local representations (one-hot)** - one dimension per object;
- **Distributed representations** - one dimension per property;
 - No single neuron encodes everything - **groups of neurons work together**;
 - **More compact and powerful**;
 - Hidden units should capture **diverse properties** of objects - not all capturing the same property - ensured by **random initialization**;
 - Initializing all units to the same weights would never break the symmetry;
 - Initializing hidden layers using **unsupervised learning** can help break the symmetry - force network to **represent latent structure** in the data; encourage hidden layers to encode **useful features**.

This can be done by using **auto-encoders**.

Auto-Encoders

Auto-encoders are feed-forward NNs trained to reproduce its input at its output layer.

- **Encoder** - maps input to a hidden representation : $h = g(Wx + b)$;
- **Decoder** - maps hidden representation to a reconstruction : $\hat{x} = W^T h(x) + c$;
- **Loss function** - $\mathcal{L}(\hat{x}, x) = \frac{1}{2} \|\hat{x} - x\|^2$;
- **Objective** - $\hat{W} = \operatorname{argmin}_W \sum_i \|W^T g(Wx_i) - x_i\|^2$.

Single Value Decomposition (SVD)

- **SVD** is a matrix factorization method that decomposes a matrix $A \in \mathbb{R}^{m \times n}$ with $m \geq n$ into the product of three matrices U , Σ and V such that $A = U\Sigma V^T$;
 - $U \in \mathbb{R}^{m \times m}$ - columns are an orthonormal basis of $R(A)$ (left singular vectors);
 - $\Sigma \in \mathbb{R}^{m \times n}$ - diagonal matrix with singular values of A ;
 - $V \in \mathbb{R}^{n \times n}$ - columns are an orthonormal basis of $R(A^T)$ (right singular vectors);

- $\sigma_1 \geq \dots \geq \sigma_r$ - square roots of the eigenvalues of $A^T A$ or AA^T
 - **singular values** of A ;
 - $U^T U = I$ and $V^T V = I$.
-

Linear Auto-Encoder

- Let $X \in \mathbb{R}^{N \times D}$ be a data matrix with N samples and D features ($N > D$);
- Assume $W \in \mathbb{R}^{K \times D}$ ($K < D$);
- We want to minimize $\sum_{i=1}^N \|x_i - \hat{x}_i\|_2^2 = \|X - XW^T W\|_F^2$;
 - $\|\cdot\|_F^2$ - **Frobenius norm**;
 - $W^T W$ has rank K ;
- From the **Eckart-Young theorem**, the minimizer is **truncated SVD** of X^T ;
 - $\hat{X}^T = U_K \Sigma_K V_K^T$;
 - $W = U_K^T$;
- This is called **Principal Component Analysis (PCA)** - fits a **linear manifold** to the data.
- By using **non-linear activations**, we obtain more sophisticated codes (representations).

There are some variants of auto-encoders:

- **Sparse auto-encoders** - add a **sparsity penalty** $\Omega(h)$ to the loss function;
 - Typically the number of hidden units is larger than the number of inputs;
 - The sparsity penalty is a **regularization** term that encourages the hidden units to be **sparse**;
- **Stochastic auto-encoders** - encoder and decoder are **not deterministic**, but involve some **noise/randomness**;
 - Uses distribution $p_{encoder}(h|x)$ for the encoder and $p_{decoder}(x|h)$ for the decoder;
 - The auto-encoder can be trained to minimize $-\log(p_{decoder}(x|h))$;
- **Denoising auto-encoders** - use a **perturbed version of the input** $\tilde{x} = x + n$, where n is a **random noise**;

- Instead of minimizing $\frac{1}{2}||\hat{x} - x||^2$, we minimize $\frac{1}{2}||\hat{x} - \tilde{x}||^2$;
- This is a form of implicit regularization that ensures **smoothness**: it forces the system to represent well not only the data points, but also their perturbations;
- **Stacked auto-encoders** - several layers of auto-encoders stacked together;
- Variational auto-encoders - learn a **latent variable model** of the data.

Regularized Auto-Encoders

- We need some sort of **regularization** to avoid **overfitting**;
- To regularize auto-encoders, **regularization** is added to the loss function;
- The goal is then to minimize $\mathcal{L}(\hat{x}, x) + \Omega(h, x)$;
- For example:
 - Regularizing the code: $\Omega(h, x) = \lambda ||h||^2$;
 - Regularizing the derivatives: $\Omega(h, x) = \lambda \sum_i ||\nabla_x h_i||^2$.

Unsupervised Pre-training

- **Unsupervised pre-training** is a technique for **initializing** the weights of a **deep NN**;
- A **greedy, layer-wise procedure**:
 - Train one layer at a time, from first to last, using **unsupervised criterion (e.g. auto-encoder)**;
 - Fix the parameters of previous hidden layers;
 - Previous layers viewed as **feature extractors**.
- After pre-training, the **whole network** is **fine-tuned** using **supervised learning - fine-tuning**;
 - Performed as in a regular **feed-forward NN** - forward propagation, backpropagation and update of weights.

Word Representations

- **Word representations** are a **key component** of **LLMs (Large Language Models)**;
- Learning representations of **words in natural language** - also called **word embeddings**;
 - An extremely successful application of representation learning;
- **Distributional similarity** - represent a word by means of its neighbors
- *you shall know a word by the company it keeps*;
- The objective is to obtain a **vector representation** for each word in a **vocabulary**; there are two main approaches:
 - Factorization of a co-occurrence word-context matrix;
 - Directly **predicting** a word from its neighbors in a **continuous word-space** - **word2vec**.

Neural Language Models

- **Embedding matrix**: assign a vector to every word in the vocabulary;
- Each word is associated with a **word embedding** - a vector of real numbers;
- Given the **context** (previous words), the **next word** is predicted;
- The word embeddings in the context window are **concatenated** into a vector that is fed to a **neural network**;
- The output of the NN is a **probability distribution** over the vocabulary
- **softmax**;
- The network is trained by a **SGD with backpropagation**.

Word2Vec

- Often, we are not concerned with language modeling, but with the **quality of the word embeddings**;
 - We do not need to predict the probability of the next word, just make sure that the true word is more likely than a random one;
- **Word2Vec** is a **shallow, two-layer NN** that is trained to **predict** the **current word** from the **context**; it comes with two variants:

- **Continuous Bag-of-Words (CBOW)** - predict the current word from the context;
- **Skip-gram** - predict the context from the current word - **more popular**.

Skip-Gram

- **Objective:** maximize the log probability of any context word given the central word:

$$J(\theta) = \frac{1}{T} \sum_{t=1}^T \sum_{-m \leq j \leq m, j \neq 0} \log p_{\theta}(x_{t+j} | x_t)$$

- There are 2 sets of parameters (2 embedding matrices - $\theta = (u, v)$):
 - Embeddings for each word o appearing as the center word - u_o ;
 - Embeddings for each word c appearing in the context of another word - v_c ;
- Uses a **log-bilinear model**: $p_{\theta}(x_{t+j} = c | x_t = o) \propto \exp(u_o^T v_c)$;
- Every word gets two vectors;
- In the end, we only care about the **word vectors** u , the **context vectors** v are discarded.

Large Vocabulary Problem

- The **softmax** is expensive to compute for large vocabularies, so there are some alternatives:
 - Stochastic sampling;
 - Noise contrastive estimation;
 - **Negative sampling**.

Negative Sampling

- **Key idea:** replace the **softmax** by **binary logistic regressions** for a true pair (**center word, context word**) and k random pairs (**center word, random word**):

$$J_t(\theta) = \log \sigma(u_o^T v_c) + \sum_{i=1}^k \log \sigma(-u_o^T v_{j_i}), j_i \sim P(x)$$

- There are several strategies for sampling the random words;
- Negative sampling is a **simple form of unsupervised pre-training**.

Linear Relationships

- **Word embeddings** are good at encoding dimensions of similarity;
- **Word analogies** can be solved well simply via subtraction in the embedding space;
- A simple way to visualize the word embeddings is to use **PCA** to project them into 2D;
- There are other methods for obtaining word embeddings:
 - **GloVe** - Global Vectors for Word Representation;
 - **FastText** - subword embeddings.