

# Deep Learning

**Deep learning** is a subfield of machine learning that is concerned with algorithms inspired by the structure and function of the brain called artificial **neural networks**.

## Recap

- [Linear Algebra](#);
- [Probability and Statistics](#);
- [Optimization](#);
- [Machine Learning](#).

## Syllabus

1. [Linear Models](#) - linear regression, perceptron, logistic regression, regularization;
2. [Neural Networks](#);
3. [Representation Learning](#);
4. [Convolutional Neural Networks](#);
5. [Recurrent Neural Networks](#)
6. [Sequence-to-Sequence Models](#) - introduction to Attention Mechanisms;
7. [Attention Mechanisms and Transformers](#);
8. [Self-Supervised Learning and Large Pretrained Models](#);
9. [Deep Generative Models](#);
10. Interpretability and Fairness.

---

The following is a summary of the course's contents.

## Linear Models

**Linear models** are a class of **regression** and **classification** models in which the prediction is a linear function of the input variables.

**Linear Regression:**  $y = w^T x + b$

- $w$  is a  $d$ -dimensional vector of **weights**;
- $b$  is a **bias** - usually included in  $w$  as a constant feature  $x_0 = 1$ ;
- Given **training data**  $D = \{(x_n, y_n)\}_{n=1}^N$ , we want to find the **best**  $w$  and  $b$ , so we use want to fit the model, i.e. find the best  $w$  and  $b$ , **minimizing the loss function** - usually, the **square loss**:  $L(w, b) = \sum_{n=1}^N (y_n - (w^T x_n + b))^2$ ;

- **Closed-form solution:**  $w = (X^T X)^{-1} X^T y$ , where  $X = \begin{bmatrix} x_1^T \\ x_2^T \\ \vdots \\ x_N^T \end{bmatrix}$  and

$$y = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_N \end{bmatrix};$$

- **Regularization** is a technique used to **reduce overfitting** by **constraining the weights** of the model;
  - **Ridge regression** is a **linear regression** model with **regularization**;
    - \*  $L(w, b) = \sum_{n=1}^N (y_n - (w^T x_n + b))^2 + \lambda \|w\|_2^2$ ;
    - \* **Closed-form solution:**  $w_{\hat{ridge}} = (X^T X + \lambda I)^{-1} X^T y$ ;

## Maximum A Posteriori (MAP) Estimation

- MAP estimation is a **Bayesian** approach to **estimation**, used to **estimate the parameters of a distribution** - used in **regularization**;
- **Bayes' rule:**  $P(\theta|x) = \frac{P(x|\theta)P(\theta)}{P(x)}$ ;
  - $\hat{w}_{MAP} = \operatorname{argmax}_w p(w|y) = \operatorname{argmax}_w p(y|w)p(w)/p(y) = \operatorname{argmin}_w \lambda \|w\|_2^2 + \sum_{n=1}^N (y_n - w^T \phi(x_n))^2$ ;
  - The **prior**  $\|w\|_2^2$  is the **regularization term** - regularizer;
  - The **likelihood**  $\sum_{n=1}^N (y_n - w^T \phi(x_n))^2$  is the **loss function**;
  - The **regularization constant** is  $\lambda = \frac{\sigma^2}{\tau^2}$ ;

- **Maximum likelihood estimation (MLE)** is a **Bayesian** approach to **estimation**;
- $\hat{w}_{MLE} = \operatorname{argmax}_w p(y|w) = \operatorname{argmin}_w \sum_{n=1}^N (y_n - w^T \phi(x_n))^2$ ;
- The **likelihood**  $\sum_{n=1}^N (y_n - w^T \phi(x_n))^2$  is the **loss function**;
- **MLE** is a **special case** of **MAP** estimation, where  $\lambda = 0$ ;

**Binary Classification -> Perceptron:**  $y = \operatorname{sign}(w^T x + b)$

- **Perceptron** is a **linear model** for **binary classification**;
- Usually, the **bias**  $b$  is included in  $w$  as a constant feature  $x_0 = 1$ , and the  $w$  vector is **augmented** with  $b$ ;
- The  $x$  vector can be represented as  $\phi(x)$ , where  $\phi$  is a **feature map**;
- Algorithm:
  1. Initialize  $w$  to zero:  $w_0 = 0$ ;
  2. While not converged, for each  $(x_n, y_n)$  in  $D$ :
    - (a) Predict:  $\hat{y} = \operatorname{sign}(w^T x_n)$ ;
    - (b) If  $\hat{y} \neq y_n$ :
      - i. Update:  $w_{t+1} = w_t + y_n x_n$ ;
      - ii. Go to step 2;
- **Perceptron convergence theorem:** if the **training data** is **linearly separable**, the **perceptron algorithm** will **converge** in a **finite number of steps**;
- It **cannot** be used for **non-linearly separable data** - XOR problem;

**Binary Classification -> Logistic Regression:**  $y = \sigma(w^T x + b)$

- **Logistic regression** is a **linear model** for **binary classification** - differs from perceptron, because it uses a **sigmoid function (continuous)** instead of a **sign function (discrete)**;
- $P_W(y|x) = \frac{\exp(w_y^T x + b_y)}{\sum_{y' \in Y} \exp(w_{y'}^T x + b_{y'})}$ , where  $W = \{w_y, b_y\}_{y \in Y}$ ;
  - Set weights to maximize conditional log-likelihood:  $L(W) = \sum_{n=1}^N \log P_W(y_n | x_n)$ ;
- No closed-form solution, so we use **stochastic gradient descent**;
- Update rule:  $w_{y_n}^{k+1} = w_{y_n}^k + \eta x_n - \eta \sum_{y' \in Y} P_{W^k}(y' | x_n) x_n$ ;

## Multi-class Classification

- **Multi-class classification** is a **classification** task with **more than two classes**, but there are several strategies to **reduce to binary classification**;
  - Parametrized by a **weight matrix**  $W \in \mathbb{R}^{d \times |Y|}$  and a **bias vector**  $b \in \mathbb{R}^{|Y|}$ ;
  - $\hat{y} = \operatorname{argmax}_{y \in Y} (W\phi(x) + b)$ ;
  - Multi-class perceptron algorithm:
    1. Initialize  $W$  to zero:  $W_0 = 0$ ;
    2. While not converged, for each  $(x_n, y_n)$  in  $D$ :
      - (a) Predict:  $\hat{y} = \operatorname{argmax}_{y \in Y} (W\phi(x_n))$ ;
      - (b) If  $\hat{y}_n \neq y_n$ :
        - i. Update:  $W_{y_n}^{k+1} = W_{y_n}^k + \phi(x_n)$
        - ii. Update:  $W_{\hat{y}_n}^{k+1} = W_{\hat{y}_n}^k - \phi(x_n)$ ;
        - iii. Go to step 2;
- 

## Neural Networks

- **Neural networks** are a class of **non-linear models** that are **inspired by the brain**;
- Consist of **neurons** that are **connected** to each other;
- **Pre-activation**:  $z(x) = w^T x + b = \sum_{i=1}^d w_i x_i + b$ ;
- **Activation**:  $h(x) = g(z(x))$ ;  $g$  can be:
  - **Linear**:  $g(z) = z$  - linear regression;
  - **Sigmoid**:  $g(z) = \frac{1}{1+e^{-z}}$  - logistic regression;
  - **Rectified Linear Unit (ReLU)**:  $g(z) = \max(0, z)$ ;
  - **Hyperbolic Tangent (tanh)**:  $g(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$ ;
  - **Softmax**:  $g(z) = \frac{e^{z_i}}{\sum_{j=1}^k e^{z_j}}$ ;
  - Others;
- A **feed-forward neural network** is a **neural network** where the **neurons** are **organized in layers** - there are **hidden layers** between the **input layer** and the **output layer**;
  - **Input layer**:  $x$  - vector of features;

- **Hidden layers:**  $h^{(1)}, h^{(2)}, \dots, h^{(L)}$ ;
  - **Output layer:**  $y$  - vector of predictions;
  - **Weights:**  $W^{(1)}, W^{(2)}, \dots, W^{(L)}$  - each weight matrix is between two layers;
  - **Biases:**  $b^{(1)}, b^{(2)}, \dots, b^{(L)}$  - each bias vector is between two layers;
  - **Hidden layer pre-activation:**  $z^{(l)}(x) = W^{(l)}h^{(l-1)}(x) + b^{(l)}$ ;
  - **Hidden layer activation:**  $h^{(l)}(x) = g(z^{(l)}(x))$ ;
  - **Output layer activation:**  $f(x) = h^{(L)}(x)$ ;
  - **Universal approximation theorem:** a feed-forward neural network with a single hidden layer can approximate any function - given enough neurons;
  - Training consists of finding the best parameters  $\theta$  - weights and biases, that minimize the loss function:  $L(\theta) := \lambda\omega(\theta) + \frac{1}{N} \sum_{n=1}^N l(f(x_n; \theta), y_n)$ ;
    - $\lambda$  is the **regularization constant**;
    - $\omega(\theta)$  is the **regularization term**;
    - $l(f(x_n; \theta), y_n)$  is the **loss function**;
    - We use **stochastic gradient descent** to minimize the loss function:  $\nabla_{\theta} L(\theta) = \lambda \nabla_{\theta} \omega(\theta) + \frac{1}{N} \sum_{n=1}^N \nabla_{\theta} l(f(x_n; \theta), y_n)$ ;
    - **Backpropagation** is a technique used to compute the gradients of the loss function with respect to the parameters - chain rule.
- 

## Representation Learning

- Neural networks can be used to learn **representations** of the data;
  - Distributed representations are exponentially more compact than one-hot representations;
  - Deeper networks exhibit hierarchical compositionality: **higher layers** represent **higher-level concepts**;
  - **Auto-encoders** are a class of **neural networks** that are used to learn **representations** of the data;
  - **Word embeddings** are a type of **representation** used to **represent words as vectors**.
-

## Convolutional Neural Networks

- **Convolutional neural networks (CNNs)** are a class of **deep neural networks** that are **specialized for processing data** that has a **grid-like topology**, such as **images**.
  - **Convolution layers** are alternated with **pooling layers** - **convolution** is a **linear operation** that **preserves the grid-like topology** of the input;
  - **Activation maps** are the **output** of a **convolutional layer**;
  - **Stride** is the **step size** of the **convolution** -  $S$ ;
  - **N of Channels** -  $K$ ;
  - **N of Filters** -  $M$ ;
  - **Padding** is the **number of zeros** added to the **input** -  $P$ ;
    - A common padding size is  $P = \frac{F-1}{2}$ , which preserves the input size;
  - Given an  $N \times N \times D$  input, a  $F \times F \times D$  filter, a stride  $S$  and padding  $P$ , the output will be a  $M \times M \times K$  activation map, where  $M = \frac{N-F+2P}{S} + 1$ ;
  - **Number of units** in a **convolutional layer**:  $M^2 \times K$ ;
  - **Number of trainable parameters** in a **convolutional layer**:  $M \times ((N^2 \times K) + bias)$ ;
  - Properties of CNNs:
    - **Invariance** - the output is **invariant** to **small translations** of the input;
    - **Locality** - the output is **only affected** by a **small region** of the input;
    - **Sparse interactions** - each output value is the result of a **small number of interactions** with the input;
    - **Parameter sharing** - the **same parameters** are used for **different parts** of the input.
- 

## Recurrent Neural Networks

- RNNs allow to take advantage of **sequential data** - words in text, DNA sequences, sound waves, etc;
  - $h_t = g(Vx_t + Uh_{t-1} + c)$  -  $h_t$  is the **hidden state** at time  $t$ ;
  - $\hat{y}_t = Wh_t + b$  -  $\hat{y}_t$  is the **output** at time  $t$ ;

- Used to generate, tag and classify sequences, and are trained using **back-propagation through time**;
    - Parameters  $V$ ,  $U$ ,  $W$ ,  $c$  and  $b$  are **shared across time steps - parameter sharing**;
  - Applications:
    - **Sequence generation** - generate a sequence of words - **auto-regressive models**;
    - **Sequence tagging** - assign a label to each element in a sequence;
    - **Pooled classification** - classify a sequence as a whole;
  - Standard RNNs suffer from vanishing and exploding gradients - alternative parameterizations like **LSTMs** and **GRUs** are used to avoid this problem;
  - **Gated Recurrent Units (GRUs)** are a type of **recurrent neural network** that are **simpler** than **LSTMs** and **perform better** than **standard RNNs** - idea is to create some **shortcuts** in the **standard RNN**;
    - $u_t = \sigma(V_u x_t + U_u h_{t-1} + b_u)$  - **update gate**;
    - $r_t = \sigma(V_r x_t + U_r h_{t-1} + b_r)$  - **reset gate**;
    - $\tilde{h}_t = \tanh(vx_t + U(r_t \odot h_{t-1}) + b)$  - **candidate hidden state**;
    - $h_t = (1 - z_t) \odot h_{t-1} + z_t \odot \tilde{h}_t$  - **hidden state**;
  - **Long Short-Term Memory (LSTM)** is a type of **recurrent neural network** that are **more complex** than **GRUs** and **perform better** than **standard RNNs** - idea is to use **memory cells**  $c_t$  to **store information**;
    - $i_t = \sigma(V_i x_t + U_i h_{t-1} + b_i)$  - **input gate**;
    - $f_t = \sigma(W_f x_t + U_f h_{t-1} + b_f)$  - **forget gate**;
    - $o_t = \sigma(W_o x_t + U_o h_{t-1} + b_o)$  - **output gate**;
    - $\tilde{c}_t = \tanh(W_c x_t + U_c h_{t-1} + b)$  - **candidate cell state**;
    - $c_t = f_t \odot c_{t-1} + i_t \odot \tilde{c}_t$  - **cell state**;
    - $h_t = o_t \odot \tanh(c_t)$  - **hidden state**.
- 

## Sequence-to-Sequence Models

- **Sequence-to-sequence models** are a class of **neural networks** that are used to **map sequences to sequences** - **encoder-decoder** architecture;
  - Used for **machine translation**, **speech recognition**, **image captioning**, etc;

- A **Neural Machine Translation (NMT)** system is a **sequence-to-sequence model** that is used to **translate** a **sequence** in one **language** to a **sequence** in another **language** - **encoder-decoder** architecture;
    - **Encoder** RNN encodes source sentence into a **vector state** -  $h_t = f(x_t, h_{t-1})$ ;
    - **Decoder** RNN decodes the **vector state** into a **target sentence** -  $y_t = g(y_{t-1}, s_t)$ ;
  - Representing the **input sequence** as a **single vector** is a **bottleneck** - **attention mechanisms** are used to **improve performance** - focus on different parts of the input.
- 

## Attention Mechanisms and Transformers

- Encoders/decoders can be RNNs, CNNs or **self-attention layers**;
  - **Self-attention** is a **linear operation** that **maps** a **sequence of vectors** to a **sequence of vectors** - **encoder-decoder** architecture;
    - **Query** vector  $q_t$ ;
    - **Key** vectors  $k_1, k_2, \dots, k_n$ ;
    - **Value** vectors  $v_1, v_2, \dots, v_n$ ;
    - **Attention weights**  $\alpha_{t,i} = \frac{\exp(q_t^T k_i)}{\sum_{j=1}^n \exp(q_t^T k_j)}$ ;
    - **Output** vector  $o_t = \sum_{i=1}^n \alpha_{t,i} v_i$ ;
  - **Transformers: encoder-decoder** architecture with **self-attention** layers instead of **RNNs**;
    - **Encoder: self-attention** layers;
    - **Decoder: self-attention** layers + **encoder-decoder attention** layers;
  - **Multi-head attention** is a **self-attention** layer with **multiple heads** - **parallel self-attention** layers;
    - **Query** vectors  $q_t$ ;
    - **Key** vectors  $k_1, k_2, \dots, k_n$ ;
    - **Value** vectors  $v_1, v_2, \dots, v_n$ ;
    - **Attention weights**  $\alpha_{t,i} = \frac{\exp(q_t^T k_i)}{\sum_{j=1}^n \exp(q_t^T k_j)}$ ;
    - **Output** vector  $o_t = \sum_{i=1}^n \alpha_{t,i} v_i$ .
-



## Self-Supervised Learning and Large Pretrained Models

- Pretraining large models and fine-tuning them to a specific task is a common practice in deep learning:
  - **Pretraining** is a technique used to **initialize** the **parameters** of a **neural network** - **self-supervised learning**;
  - **Fine-tuning** is a technique used to **adapt** the **parameters** of a **neural network** to a **specific task**;
- Models: ELMo, BERT, GPT, etc;
- **Adapters** and **prompting** are other strategies more parameter-efficient than fine-tuning;
- Current models exhibit **few-shot learning** capabilities - can be trained with **few examples**.